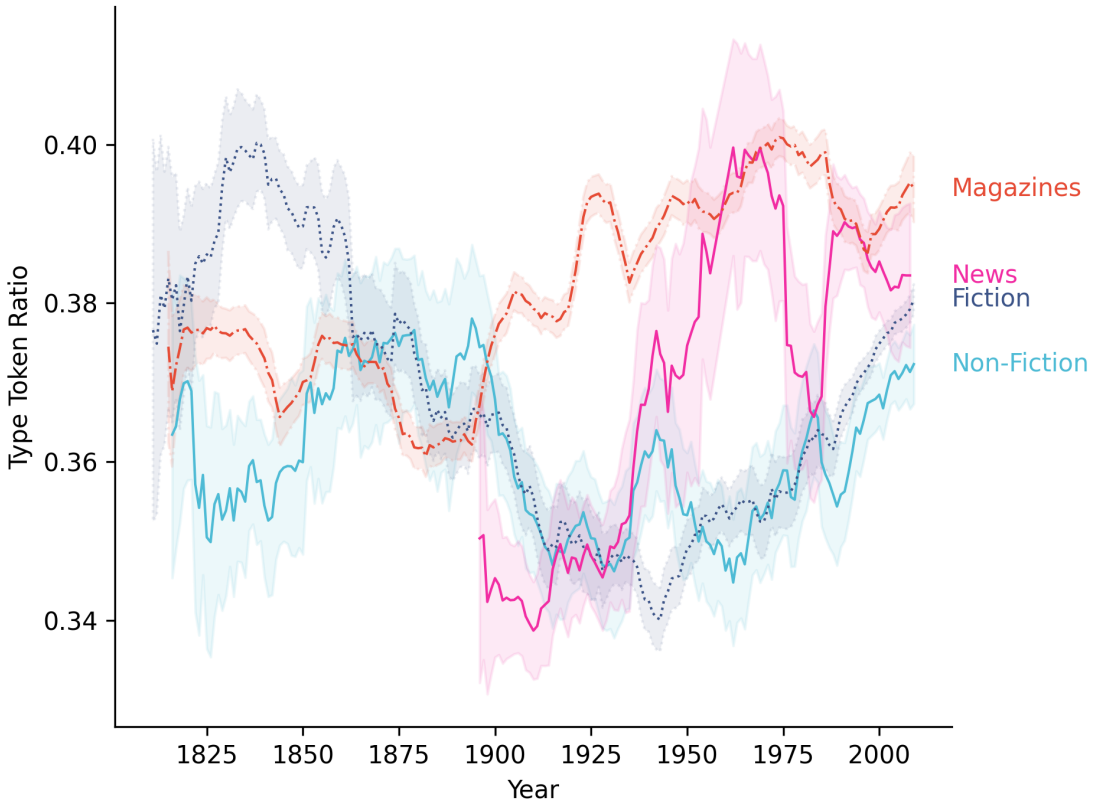

Figure 8: Historical timeseries of type token ratio in the Corpus of Historical American English. Type token ratio was calculated for text samples from COHA truncated with $N = 2000$ words. For each media category and year, a moving average of all valid samples with ± 5 years was calculated. The shaded region shows a 95% confidence interval for this average.

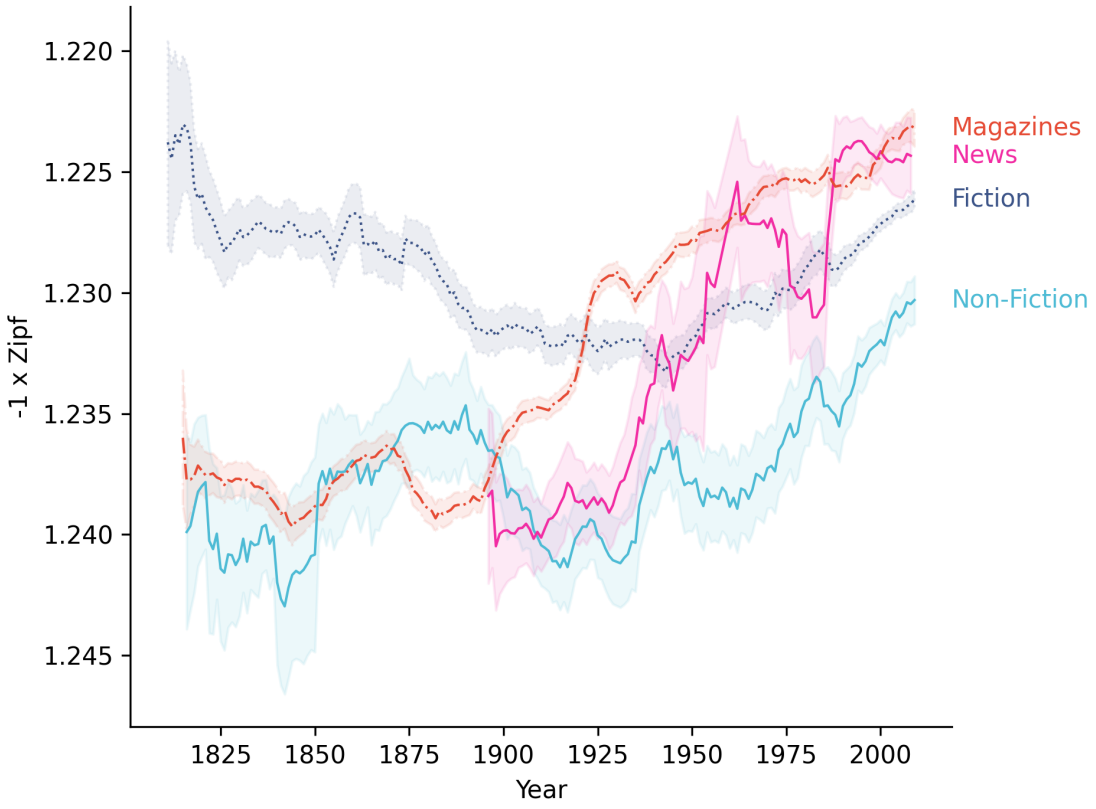

Figure 9: Historical timeseries of Zipf exponent in text samples in written media categories in American English. The timeseries was calculated in the same way as in the previous figure.

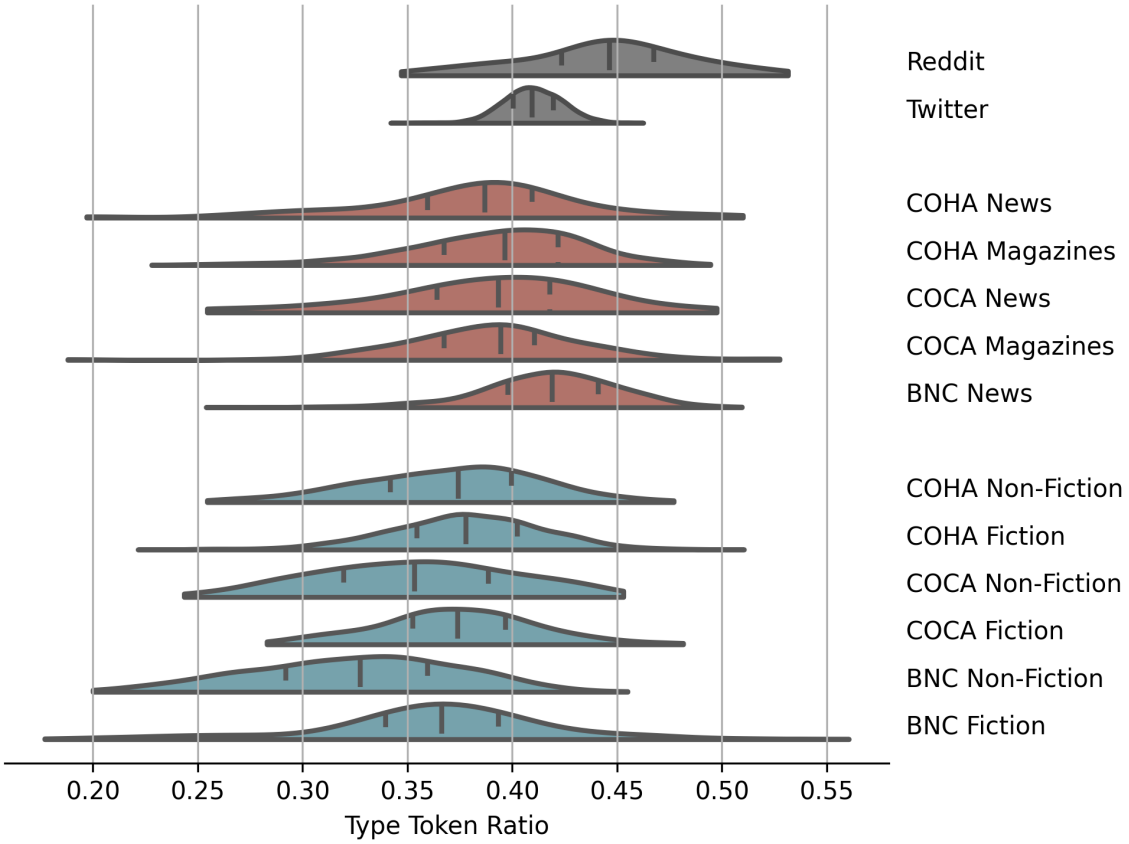

Figure 10: Distribution snapshots of type token ratio across different text corpora for text samples with $N = 2000$ words. COHA samples are from the year 2000 onwards only. Social media text samples were collated from status updates.