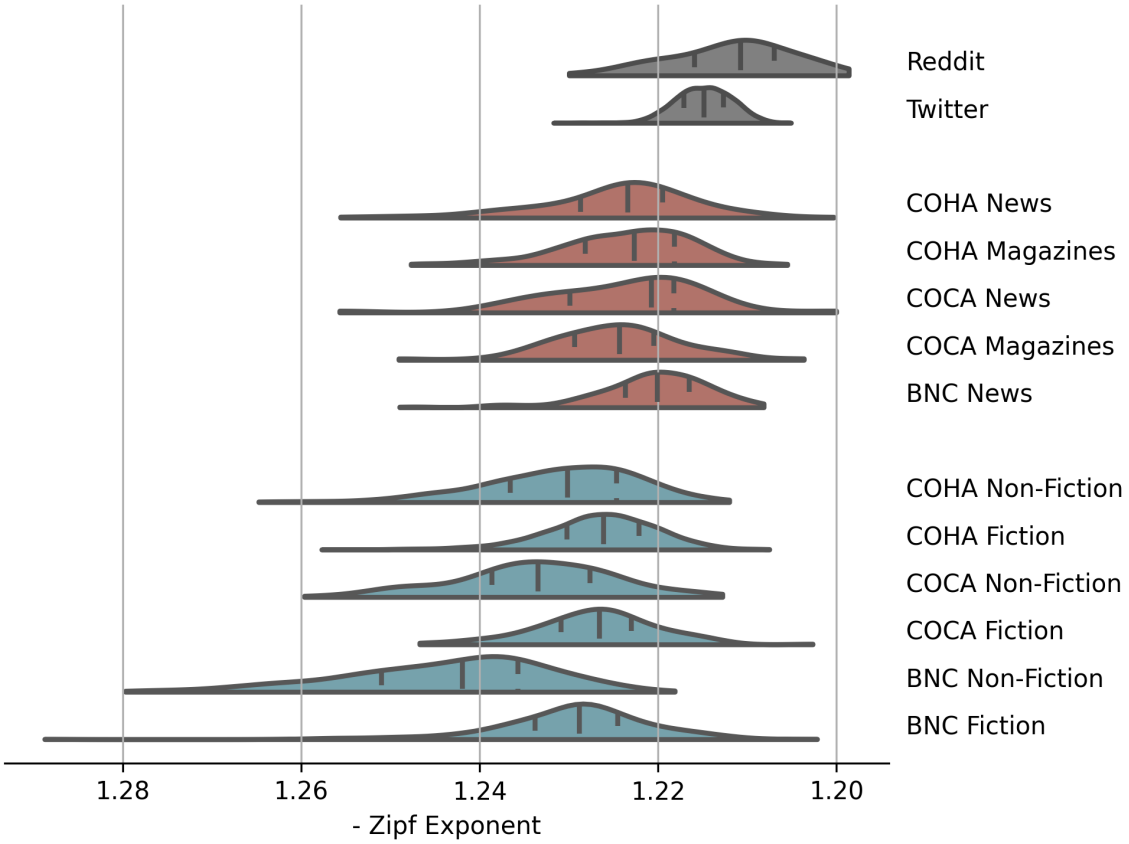

Figure 5: Distribution snapshots of the Zipf exponent across different text corpora for text samples with $N = 2000$ words. COHA samples are from the year 2000 onwards only. Social media text samples were collated from status updates.

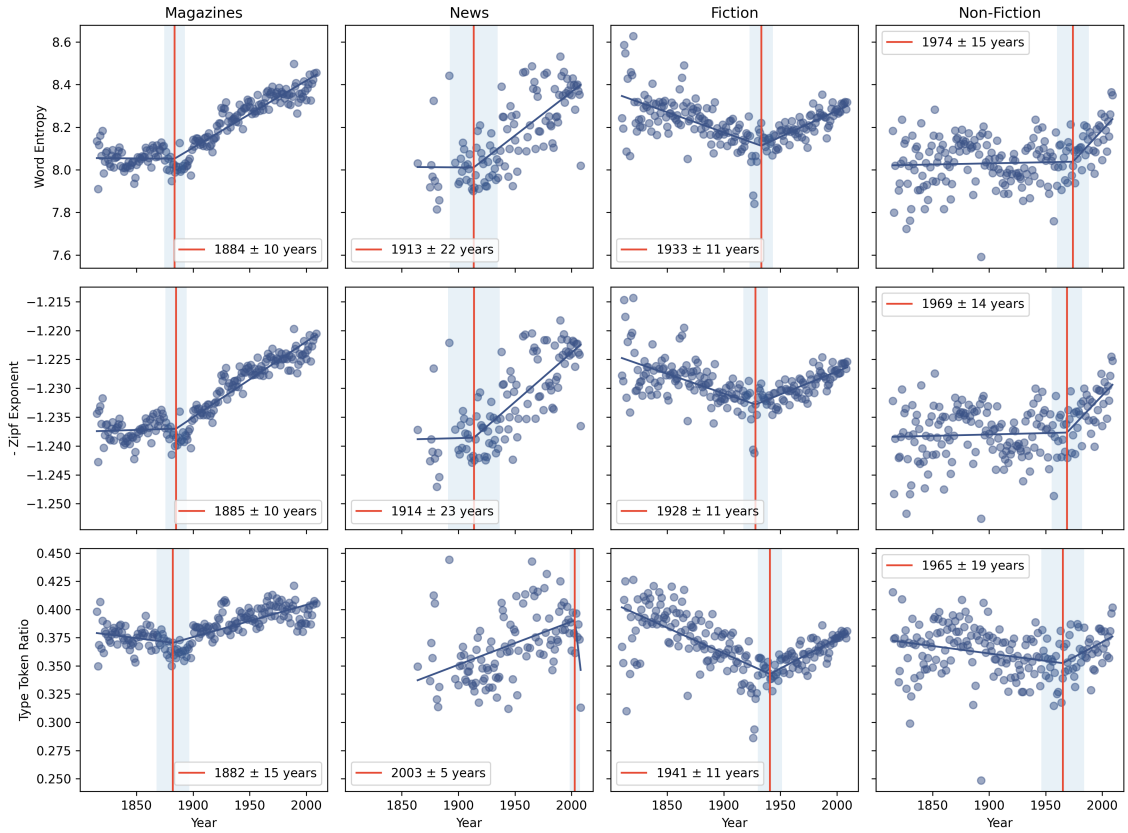

Figure 6: Median annual values for each category and lexical measure. The points were fit with a piecewise-regression, with red lines showing the estimated breakpoints. The shaded region shows a 95% confidence interval for those breakpoints.

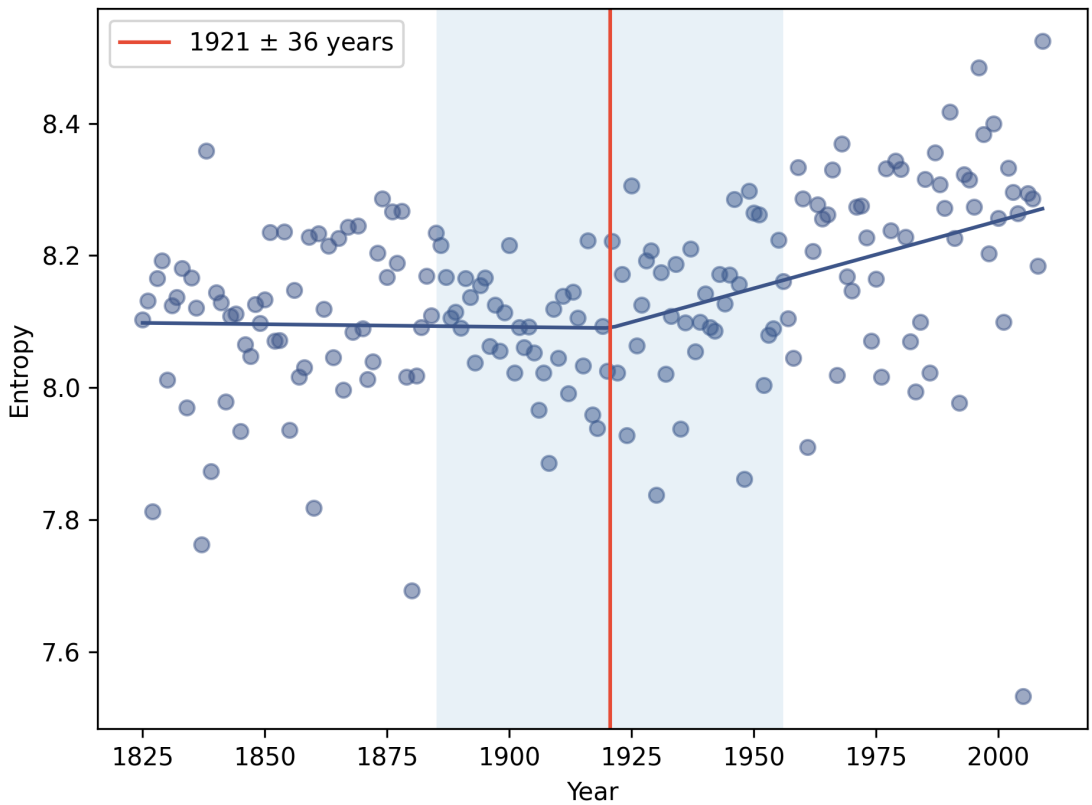

Figure 7: Mean annual values for the media categories combined for word entropy. Annual means were first found within each media category, and then averaged over the media categories. The points were fit with a piecewise-regression, with red lines showing the estimated breakpoint. The shaded region shows a 95% confidence interval for that breakpoint.